


Chapter 2

Application Layer – part 2

School of Computing
Gachon Univ.
Joon Yoo

Many slides from J.F Kurose and K.W. Ross

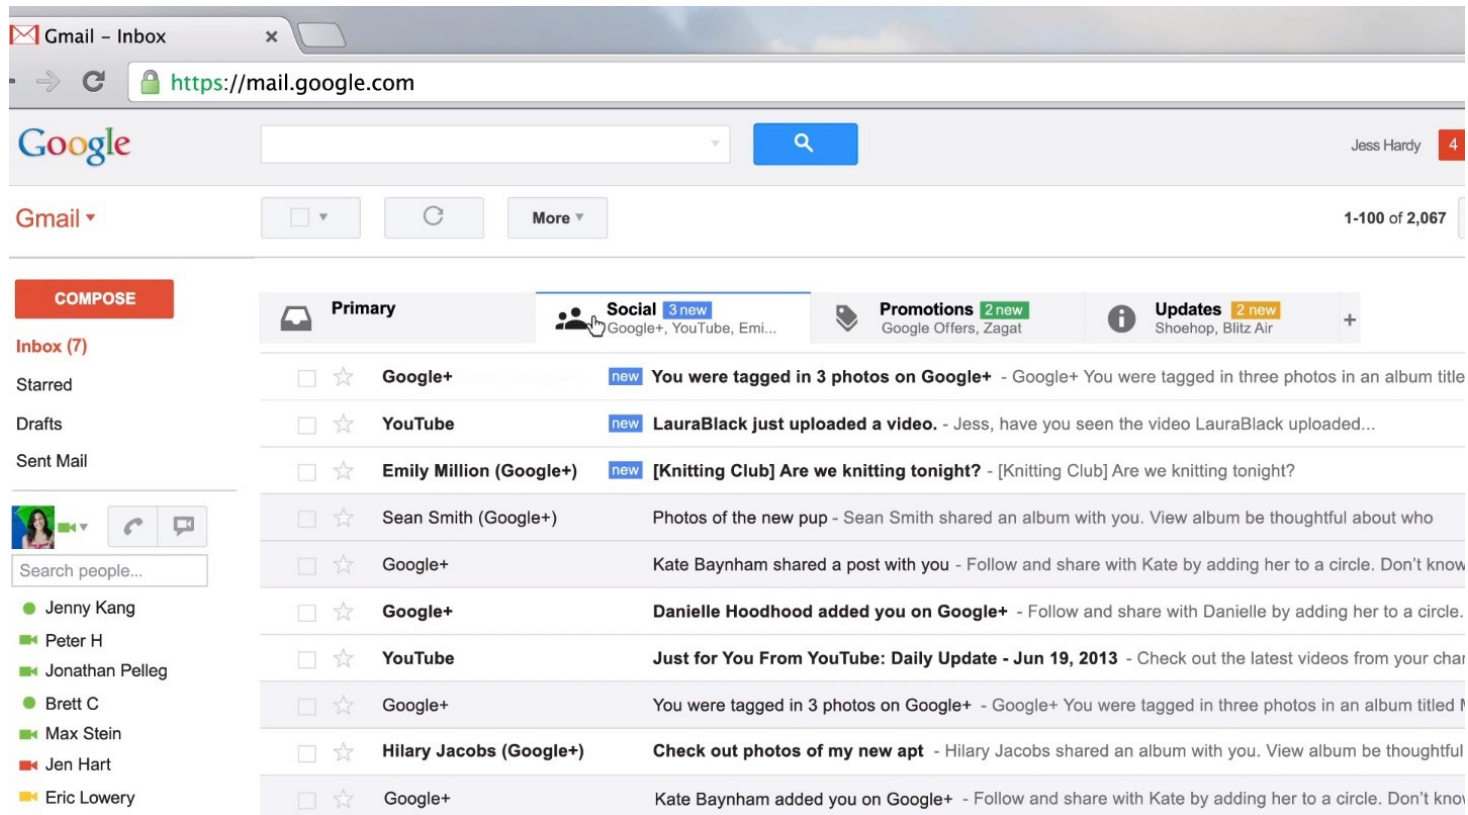


*Computer Networking:
A Top-Down Approach*

8th edition (Global edition)
Jim Kurose, Keith Ross
Pearson, 2021

Question

How can you receive your E-mail?



Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.6 Video Streaming and CDN

2.7 Socket Programming

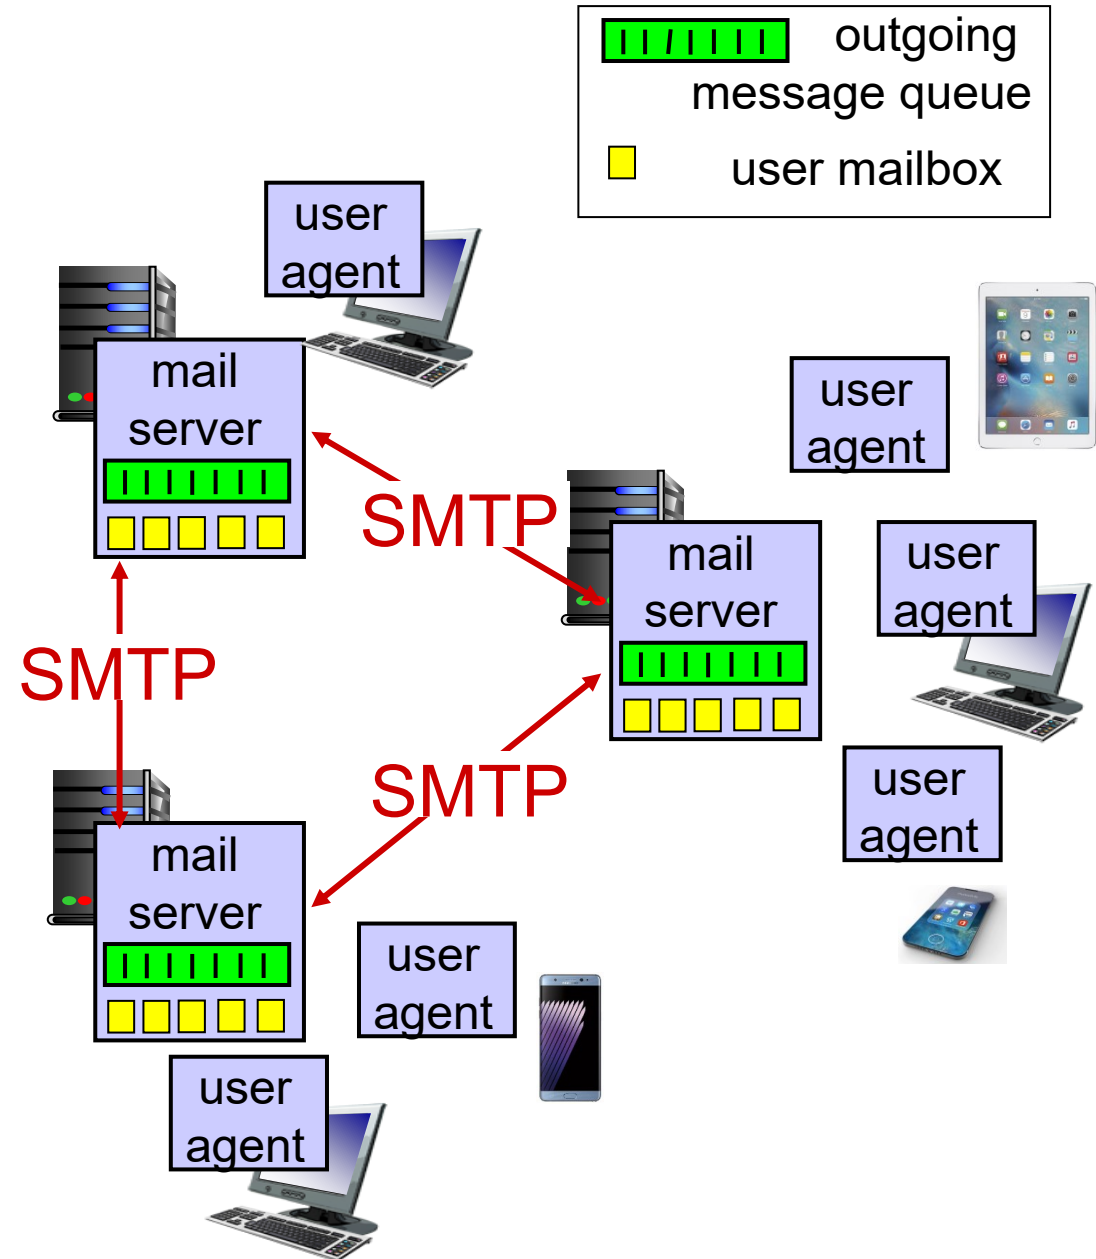
Electronic Mail

- ❖ **Electronic Mail (E-mail)** has been around since the beginning of the Internet – and remains one of the most important and utilized application
- ❖ E-mail is an *asynchronous* communication
 - People can send out messages when it is convenient for them, no need to coordinate other people's schedules
 - c.f., HTTP is *synchronous*: server-client must immediately communicate

Electronic mail

Three major components:

- ❖ user agents (UA)
- ❖ mail servers
- ❖ simple mail transfer protocol (SMTP) – application layer protocol

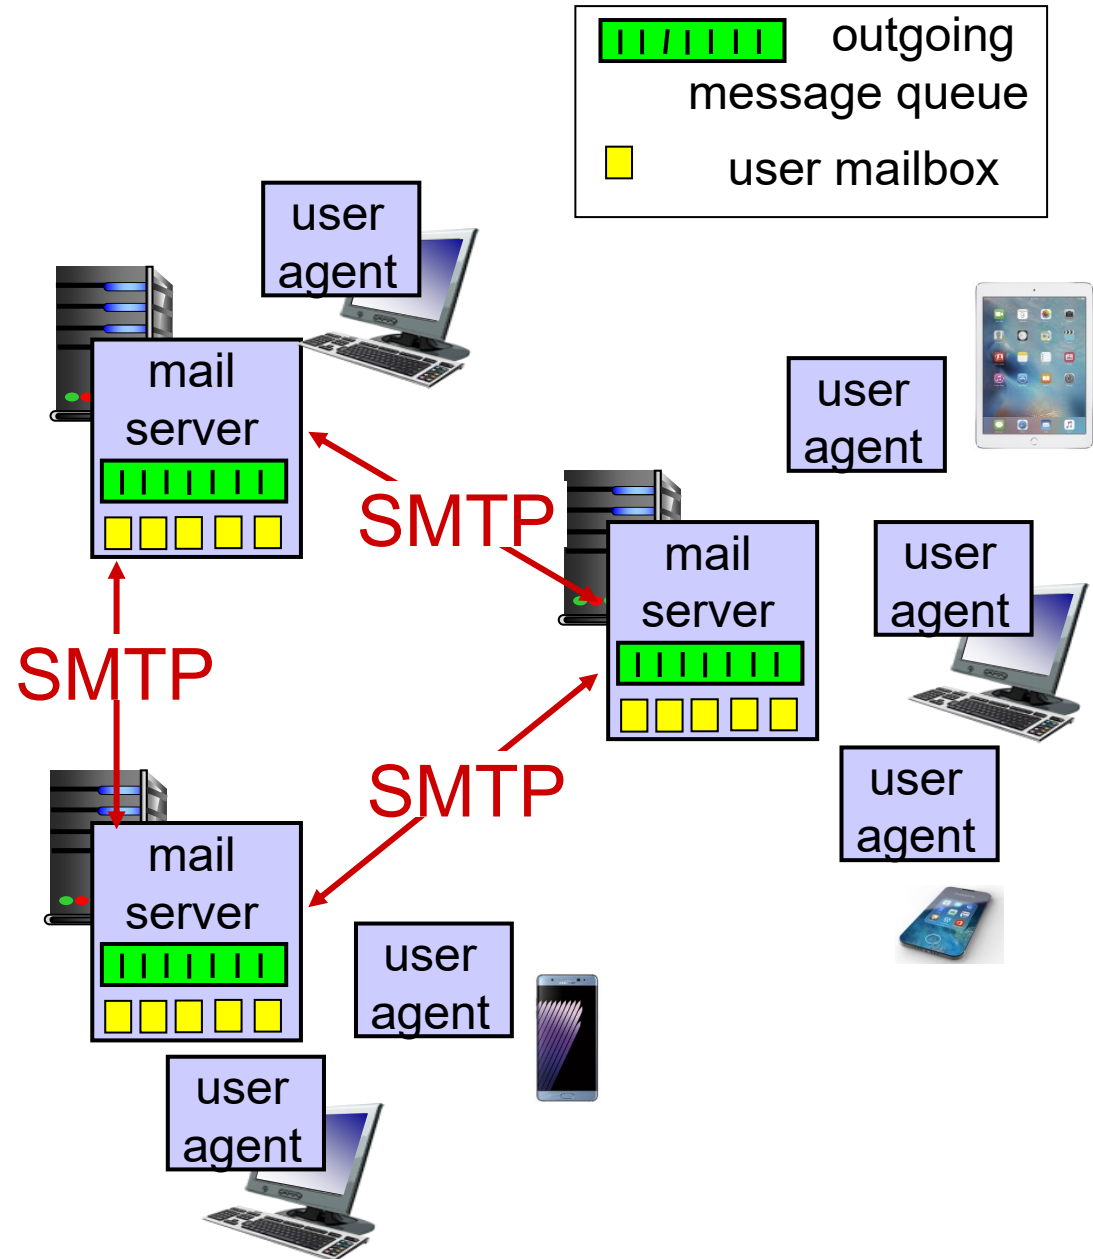
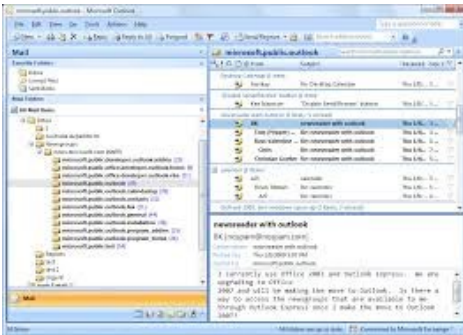




Electronic mail

User Agent (UA)

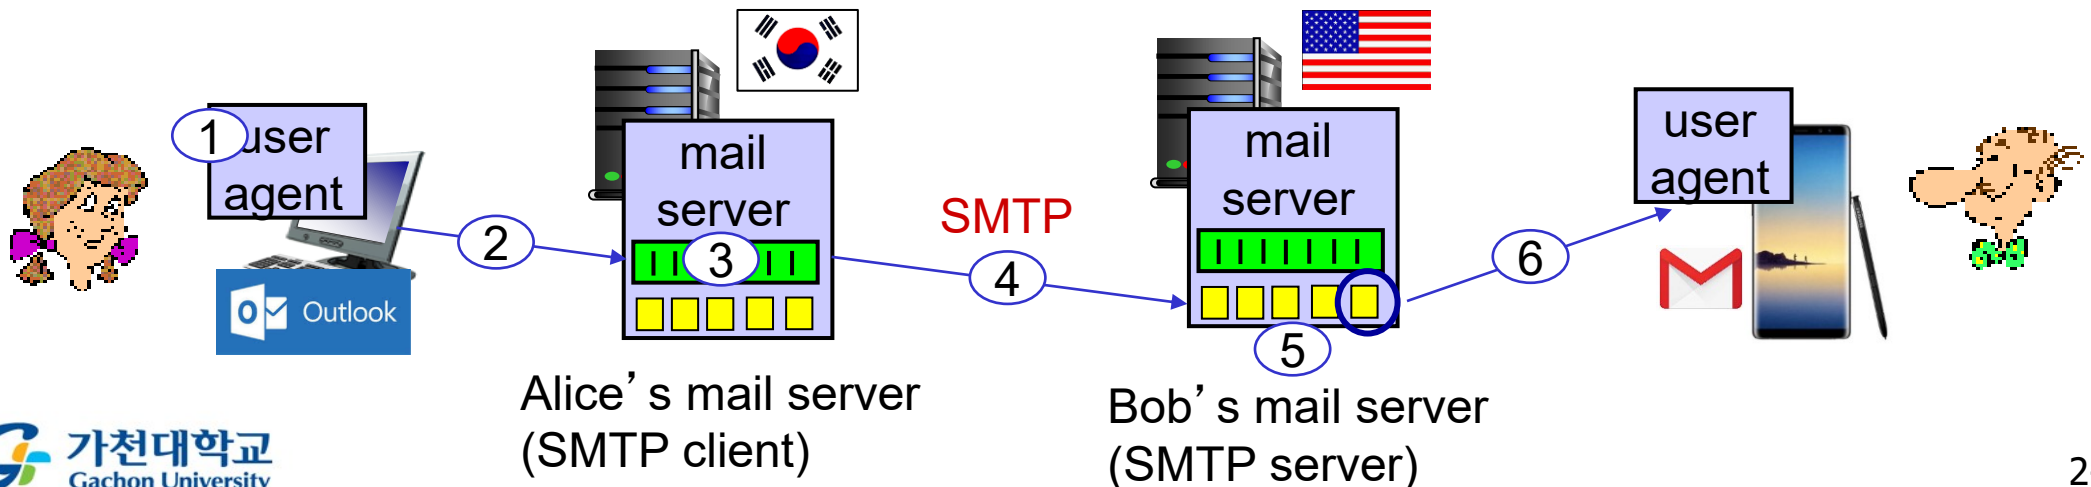
❖ allows users to read, reply to, forward, save, and compose messages

- Android Gmail client
- iPhone mail client
- MS Outlook



Scenario: Alice sends message to Bob

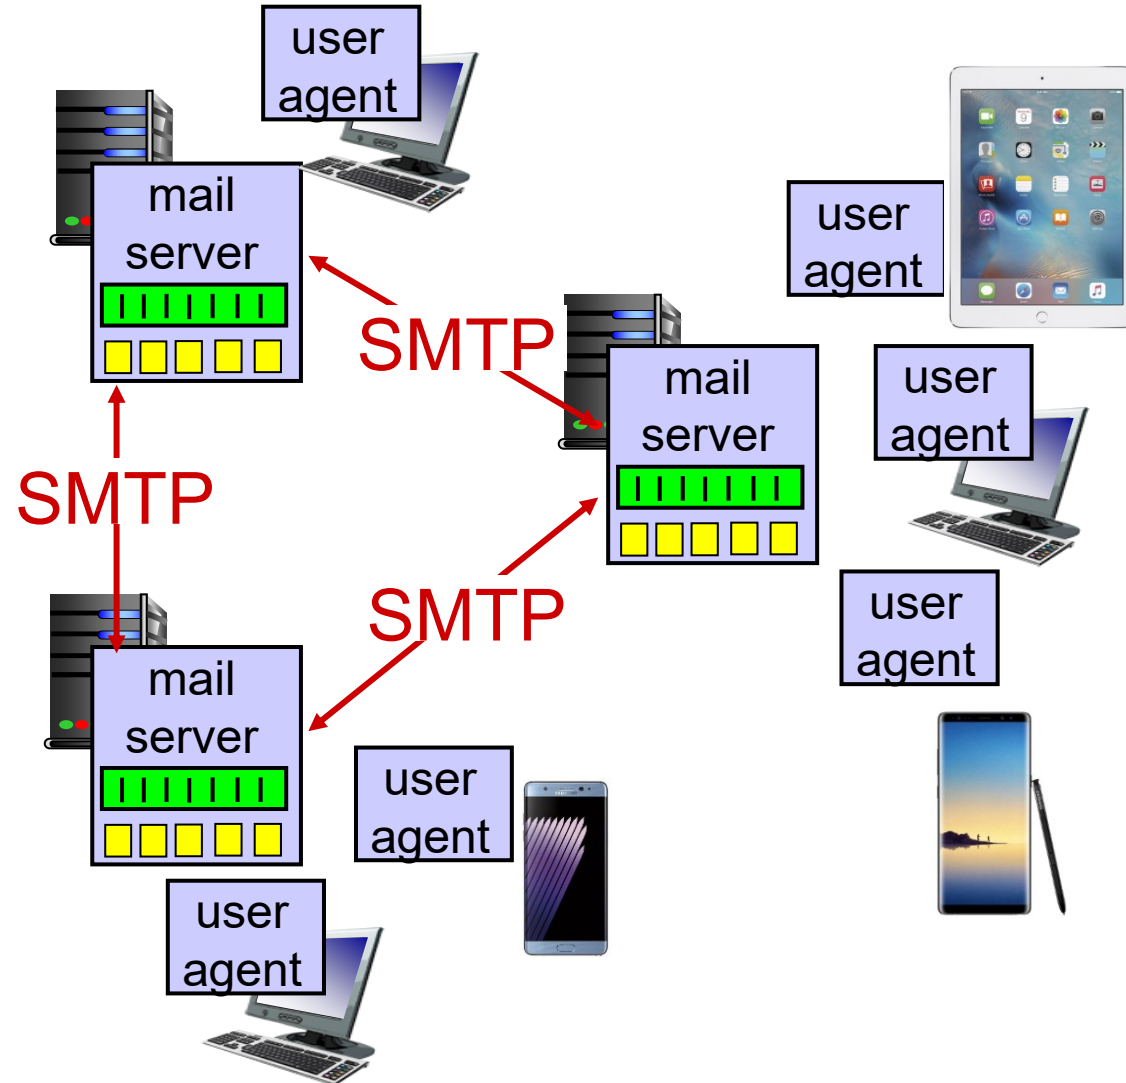
- 1) Alice uses **UA** to compose e-mail message “to” bob@some school .edu
- 2) Alice’s **UA** sends message to her **mail server**; message placed in message queue
- 3) client side of **SMTP** opens TCP connection with Bob’s **mail server**
- 4) **SMTP** client sends Alice’s message over the TCP connection
- 5) Bob’s **mail server** places the message in Bob’s mailbox
- 6) When Bob wants to read a message, his **UA** retrieves message from his mailbox



Electronic mail: mail servers


mail servers:

- ❖ *message queue* of outgoing (to be sent) mail messages
- ❖ *mailbox* contains incoming messages for user
- ❖ *SMTP protocol* between mail servers to send/receive email messages
 - Simple Mail Transfer Protocol (SMTP)



Mail Server-to-server : SMTP

- ❖ direct transfer: sending server (SMTP client) to receiving server (SMTP server)
- ❖ uses _____ transport protocol to reliably transfer email message from client to server, port 25
 - If the server is down, the client tries again later



SMTP: Comparison with HTTP

- ❖ HTTP: transfer files from Web server to Web client (browser)
 - **Pull protocol**: HTTP client pulls the information from server
 - TCP connection is initiated by machine that wants to receive
- ❖ SMTP: transfer files from one mail server to another mail server
 - **Push protocol**: sending mail server pushes the file to the receiving mail server
 - TCP connection is initiated by machine that wants to send
- ❖ SMTP requires message (header & body) to be in **7-bit ASCII**s
 - Need to encode all binary multimedia data into ASCII before sending over SMTP (No such restriction in HTTP)
 - Image (sender) → 7-bit ASCII text (in SMTP msg) → Image (receiver)
 - This made sense in early 80s when transmission capacity was scarce, so all messages were text – but now it is archaic

Try SMTP interaction for yourself:

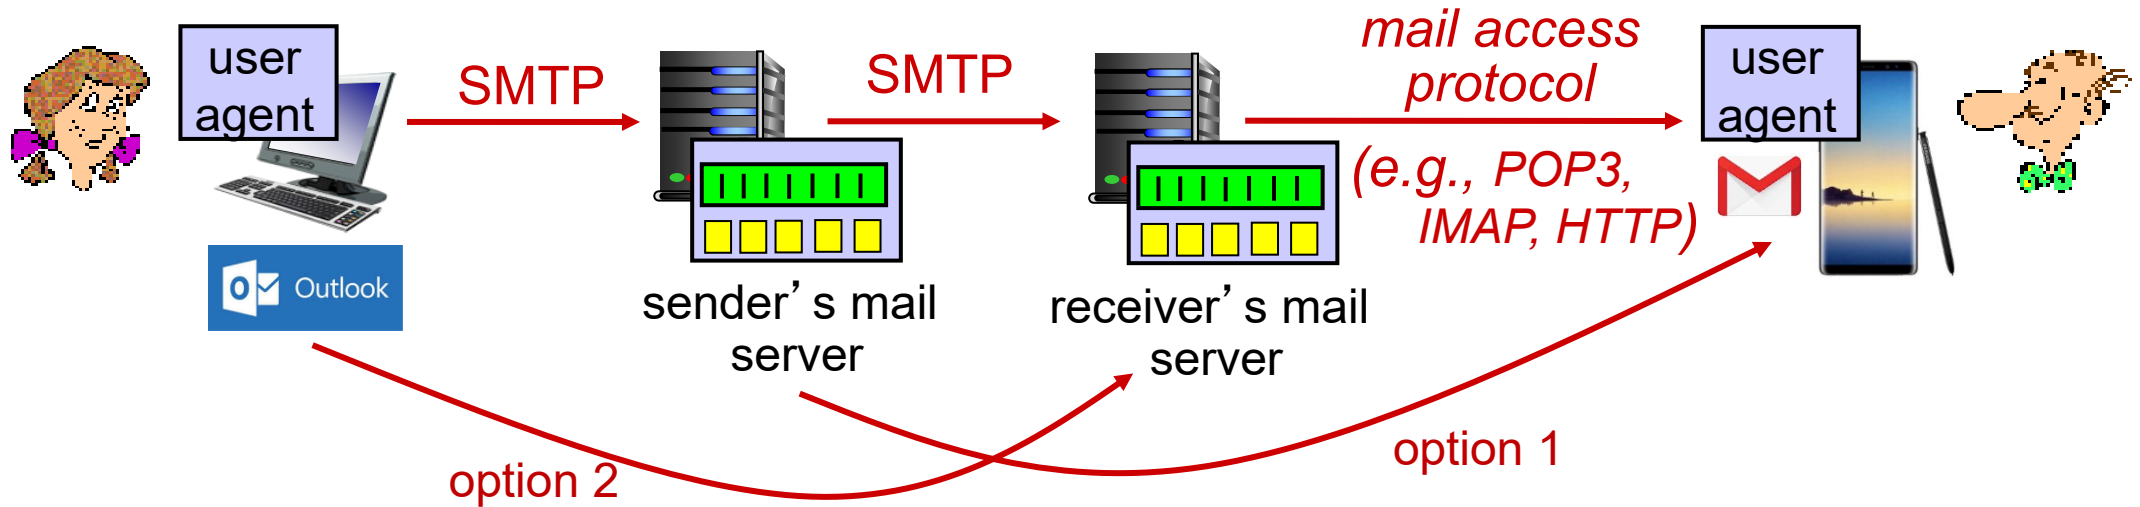
telnet <servername> 25

- see 220 reply from server
- enter HELO, MAIL FROM:, RCPT TO:, DATA, QUIT commands

above lets you send email without using e-mail client (reader)

Note: this will only work if <servername> allows telnet connections to port 25 (this is becoming increasingly rare because of security concerns)

Mail access protocols



- ❖ 3 step procedure: UA → mail server → mail server → UA
- ❖ Why not take option 1?
- ❖ Why not take option 2?
- ❖ How does Bob's UA retrieve mail from mailbox?

POP3, IMAP & Web mail

POP3 - Post Office Protocol [RFC 1939]

- ❖ Transfer mail from recipient's mail server to user agent (client)
- ❖ POP3 uses “download and delete” mode

IMAP - Internet Mail Access Protocol [RFC 1730]

- ❖ keeps all messages in one place: at server – doesn't delete

Web-based E-mail

- ❖ User agent is Web browser and communicate with mailbox via HTTP (rather than SMTP, POP3, or IMAP)
- ❖ The mail server still uses SMTP to send/receive messages to/from other mail servers
- ❖ e.g., Gmail, NAVER mail, Gachon E-mail

Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.6 Video Streaming and CDN

2.7 Socket Programming

Domain name



But, Google server's actual address is an **IP address**!

```
C:\Users\jyoo>ping google.co.kr  
Ping google.co.kr [59.18.46.113] 32바이트 데이터 사용:
```

DNS: domain name system

people: many identifiers:


- SSN (≈주민번호), driver's license#, passport #, ...

Internet hosts, routers:

- **IP address** or **host name** (e.g., www.yahoo.com)
- People prefer ____ and routers prefer ____

Q: how to map between IP address and name, and vice versa ?

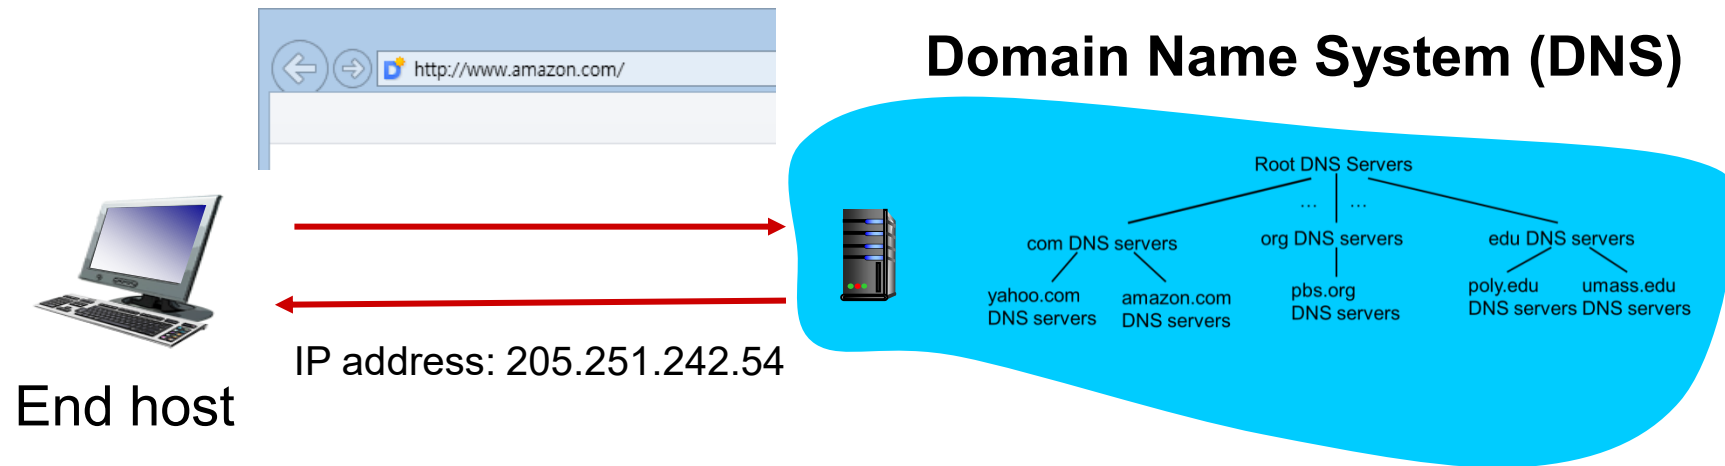
A: DNS



DNS: domain name system

Domain Name System (DNS):

- ❖ *distributed database* implemented in hierarchy of many *name servers*
- ❖ *application-layer protocol*: hosts, name servers employ DNS to *translate* host names into IP addresses



DNS: distributed vs. central

Single centralized server


- ❖ Ask one DNS server to translate name to IP address
- ❖ Very simple!

Then, why not single centralize DNS server?

- ❖ single point of failure
- ❖ traffic volume
- ❖ distant centralized database
- ❖ maintenance

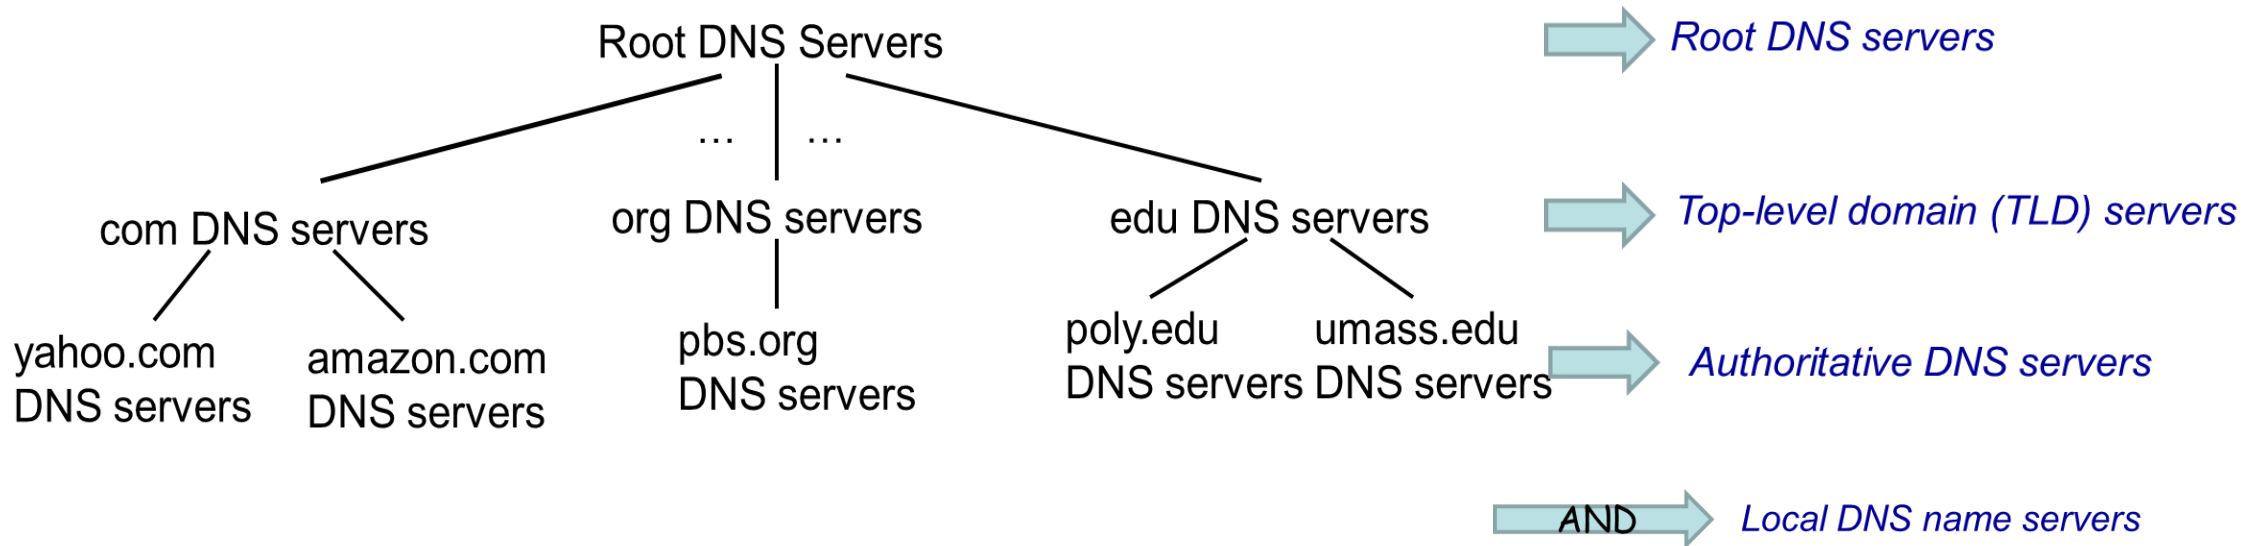


VS.



A: doesn't scale!

DNS: a distributed, hierarchical database



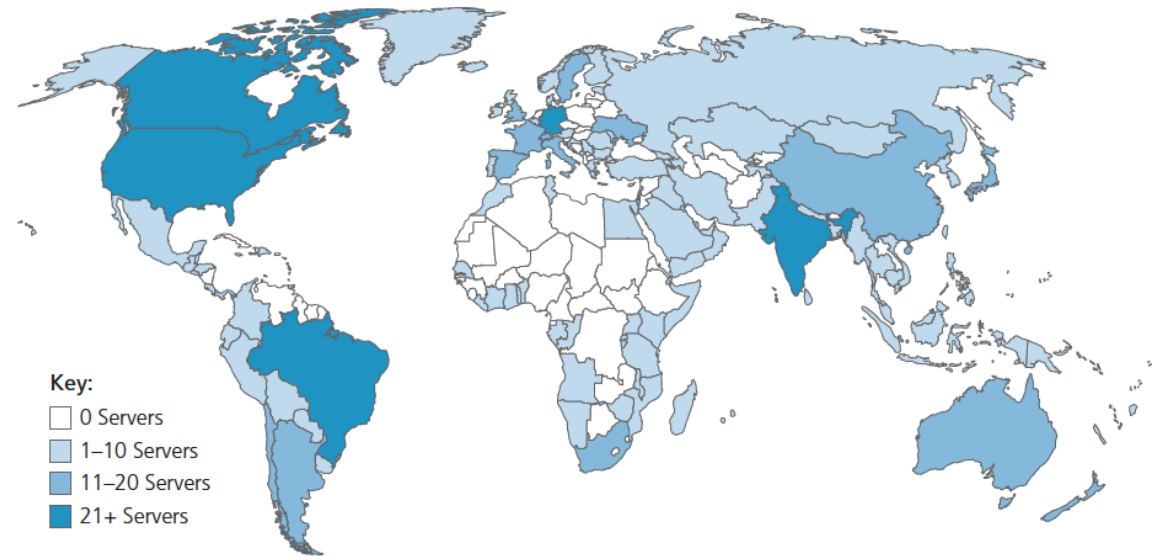
client wants IP address for `www.amazon.com` ; 1st approx:

- ❖ client queries **root server** to find TLD (`.com`) DNS server
- ❖ client queries `.com` DNS server to get authoritative (`amazon.com`) DNS server
- ❖ client queries `amazon.com` DNS server to get IP address for `www.amazon.com`

DNS: root Name Servers

- official, contact-of-last-resort by name servers that can not resolve name
- *incredibly important* Internet function
 - Internet couldn't function without it!
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain

13 logical root name “servers”
worldwide each “server” replicated
many times (~200 servers in US)



TLD, authoritative servers

Top-Level Domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: kr, uk, fr, ca, jp

Authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider
- e.g., google, NAVER, Gachon

Local DNS name server

- ❖ each ISP (residential ISP, company, university) has one (or more)

- also called “default name server”

- ❖ when host makes DNS query, query is sent to its local DNS server

- has local cache of recent name-to-address translation pairs (but may be out of date!)
- acts as proxy, forwards query into hierarchy

ipconfig/all

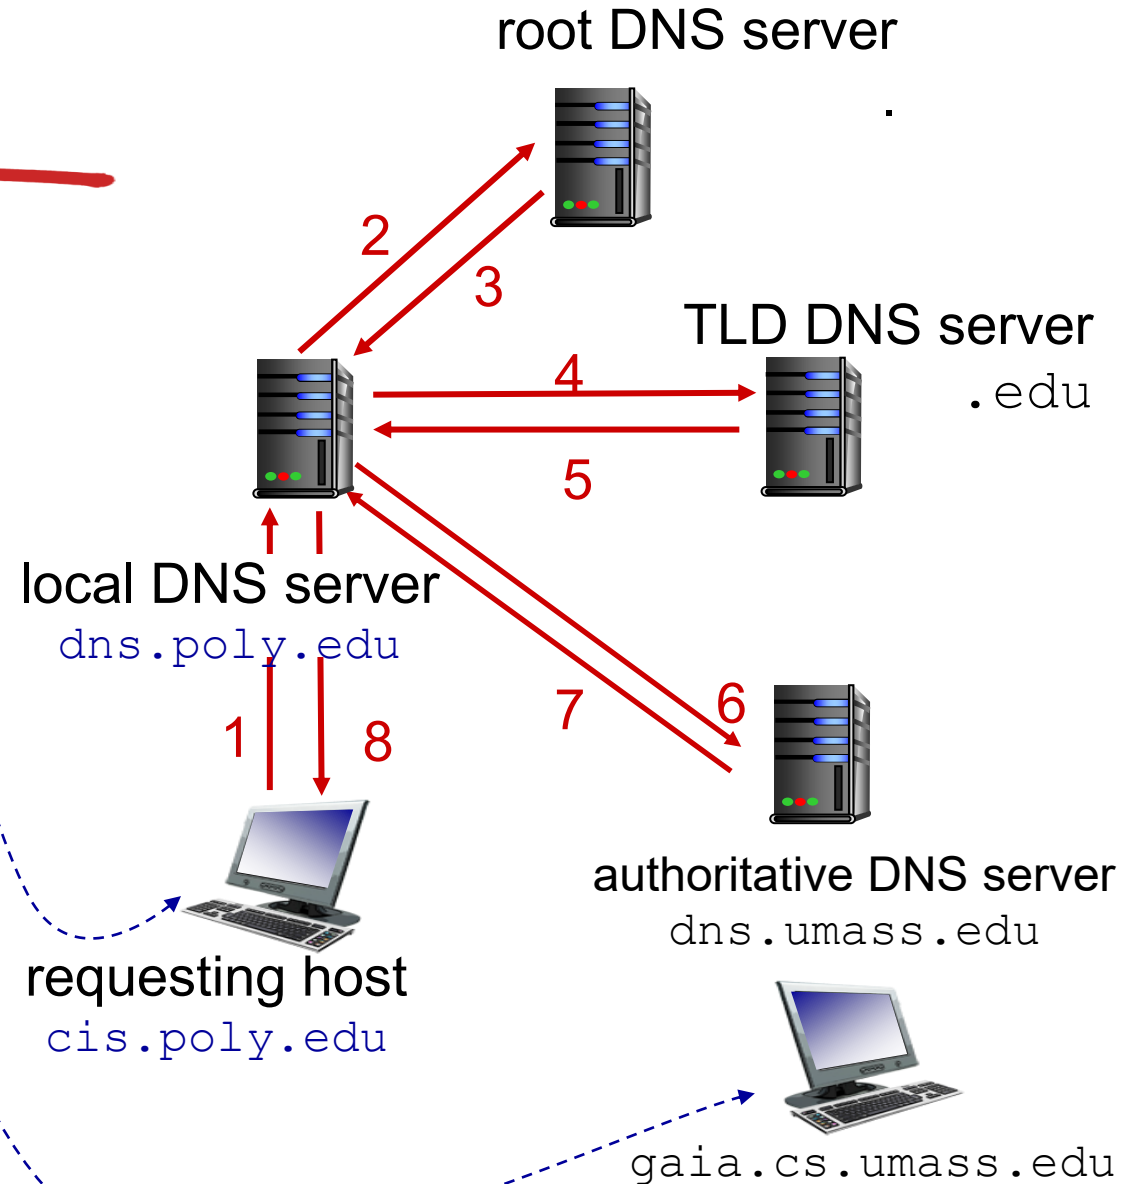
```
IPv4 주소 . . . . . : 121.135.107.150<기본 설정>
서브넷 마스크 . . . . . : 255.255.255.0
임대 시작 날짜 . . . . . : 2014년 9월 21일 일요일 오후 1:59:47
임대 만료 날짜 . . . . . : 2014년 9월 21일 일요일 오후 11:11:20
기본 게이트웨이 . . . . . : 121.135.107.1
DHCP 서버 . . . . . : 121.137.7.58
DHCPv6 IAID . . . . . : 199754650
DHCPv6 클라이언트 DUID. . . . . : 00-01-00-01-18-39-90-22-E8-03-9A-66-87-44
DNS 서버 . . . . . : 168.126.63.1
                  168.126.63.2
Tcpip를 통한 NetBIOS. . . . . : 사용
```

DNS name resolution example

- ❖ host at `cis.poly.edu` wants IP address for `gaia.cs.umass.edu`

iterated query:

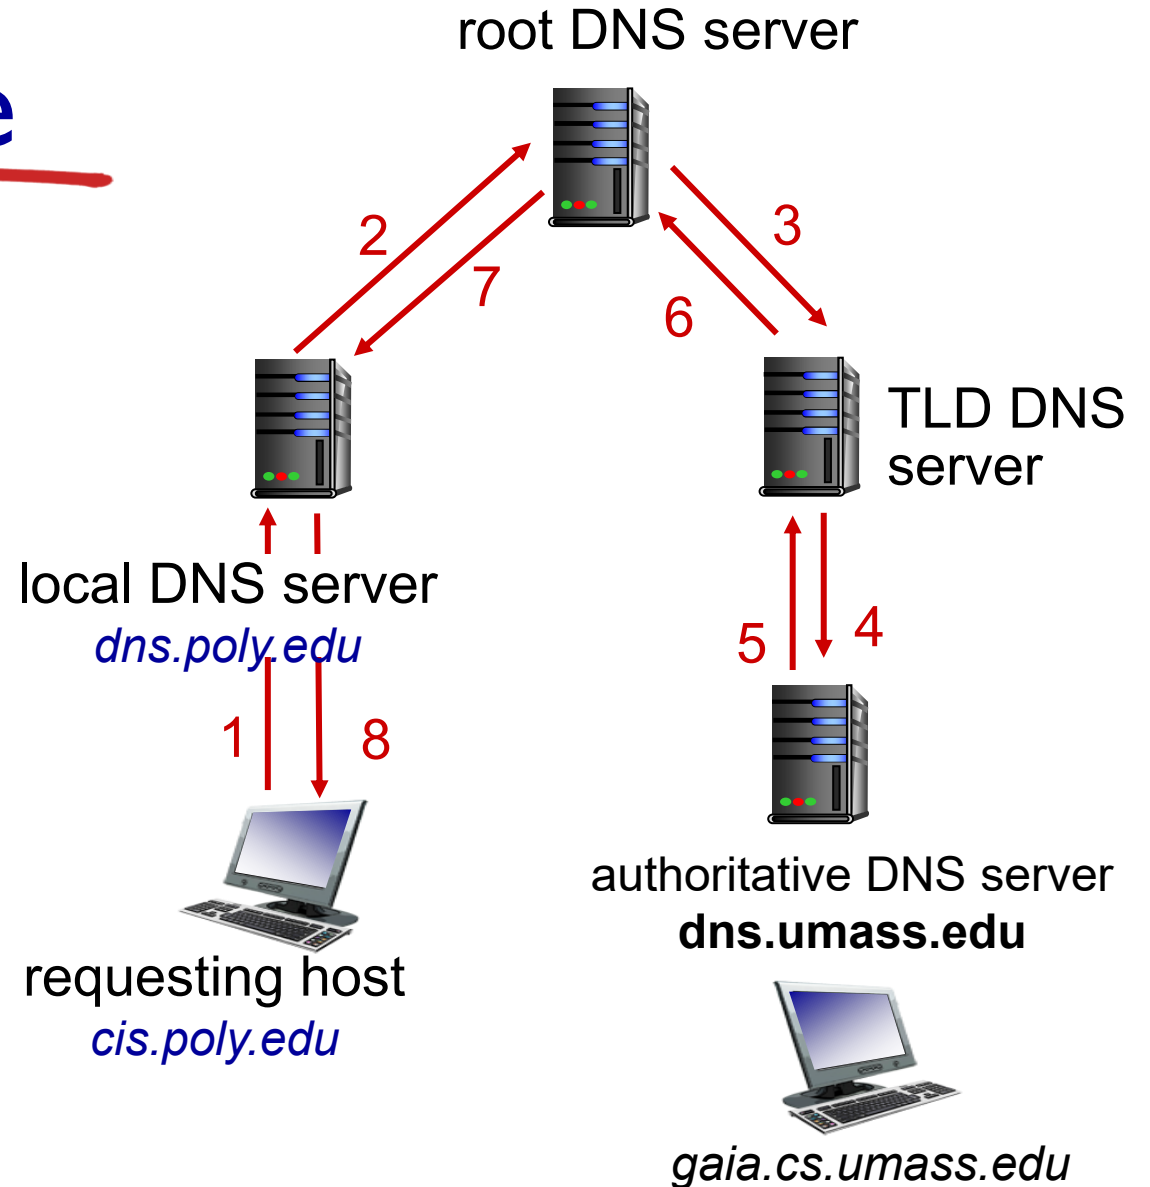
- ❖ contacted server replies with name of server to contact
- ❖ “I don’t know this name, but ask this server”



DNS name resolution example

recursive query:

- ❖ puts burden of name resolution on contacted name server
- ❖ heavy load at upper levels of hierarchy?



DNS: caching, updating records

- ❖ once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time-to-live (TTL)
 - TLD servers typically cached in local name servers
 - thus root name servers not often visited
- ❖ cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire

Question

How can you watch a YouTube video without buffering?



Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.6 Video Streaming and CDN

2.7 Socket Programming

Video Streaming and CDNs: context


- stream video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube, Amazon Prime: 80% of residential ISP traffic (2020)
- challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
- challenge: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution: distributed, application-level infrastructure*



Multimedia: video

- video: sequence of images displayed at constant rate
 - e.g., 24 images/sec (or fps)
- digital image: array of pixels
 - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i




frame $i+1$

Multimedia: video

■ examples:

- MPEG I (CD-ROM) 1.5 Mbps
- MPEG2 (DVD) 3-6 Mbps
- MPEG4 (often used in Internet, 64Kbps – 12 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i


temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i



frame $i+1$

Streaming stored Video

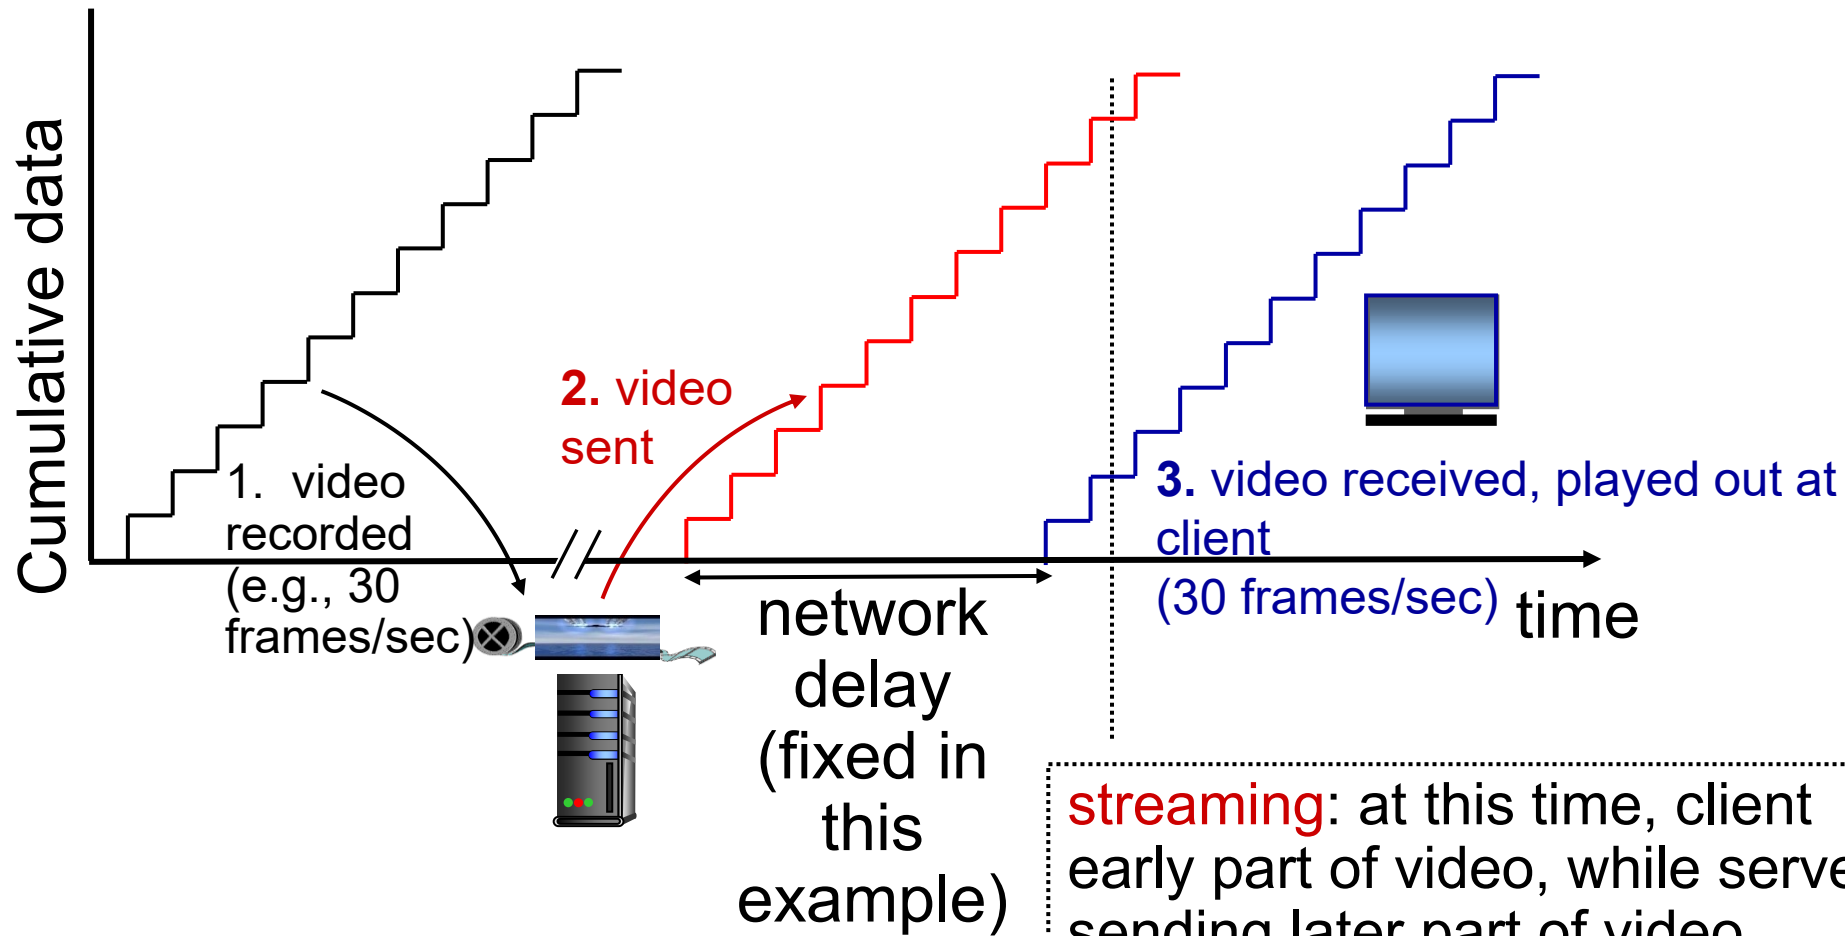
simple scenario:



Main challenges:

- ❖ server-to-client bandwidth will *vary* over time, with changing network congestion levels (in house, in access network, in network core, at video server)
- ❖ packet loss and delay due to congestion will delay playout, or result in poor video quality


Streaming stored Video



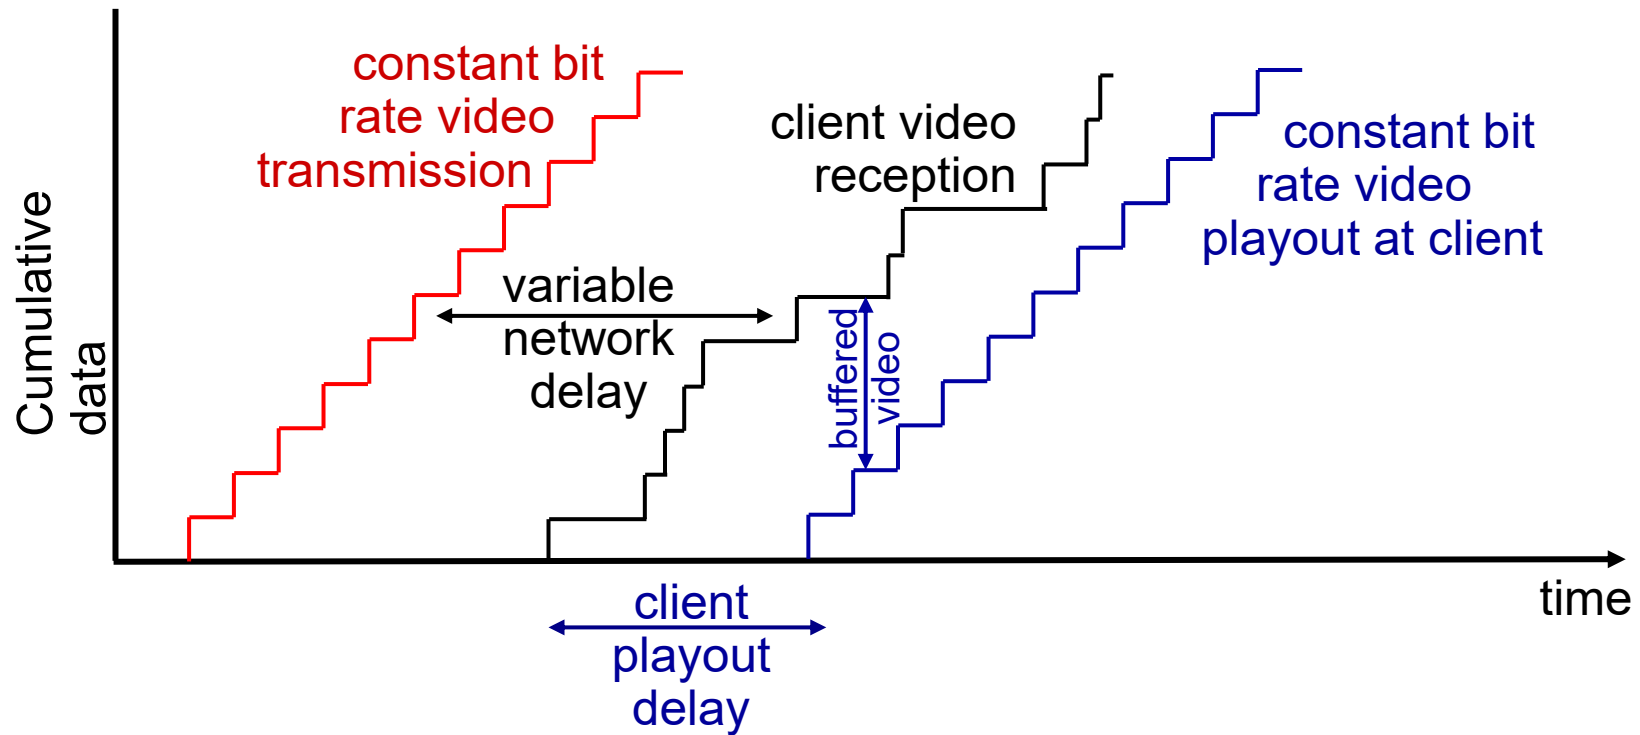
streaming: at this time, client playing out early part of video, while server still sending later part of video

Streaming stored Video: challenges

- **continuous playout constraint**: once client playout begins, playback must match original timing
 - ... but **network delays are variable** (jitter), so will need **client-side buffer** to match playout requirements
- other challenges:
 - client interactivity: pause, fast-forward, rewind, jump through video
 - video packets may be lost, retransmitted



Streaming stored Video: playout buffering



- *client-side buffering and playout delay*: compensate for network-added delay, delay jitter

Streaming multimedia: DASH


- **DASH**: *D*ynamic, *A*daptive *S*treaming over *H*TTP

- **server**:

- divides video file into multiple chunks
- each chunk stored, encoded at different rates
- *manifest file*: provides URLs for different chunks

- **client**:


- periodically measures server-to-client bandwidth
- consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)



Streaming multimedia: DASH

- “intelligence” at client: client determines

- *when* to request chunk (so that buffer starvation, or overflow does not occur)
- *what encoding rate* to request (higher quality when more bandwidth available)
- *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)



Streaming video = encoding + DASH + playout buffering

Content distribution networks (CDNs)

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **option 1:** single, large “mega-server”
 - single point of failure
 - point of network congestion
 - long path to distant clients
 - multiple copies of video sent over outgoing link

....quite simply: this solution **doesn't scale**


Content distribution networks (CDNs)

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **option 2:** store/serve multiple copies of videos at multiple geographically distributed sites (**CDN**)
 - **enter deep:** push CDN servers deep into many access networks
 - close to users
 - Akamai: 240,000 servers deployed in more than 120 countries (2015)
 - **bring home:** smaller number (10's) of larger clusters near (but not within) access networks
 - used by Limelight

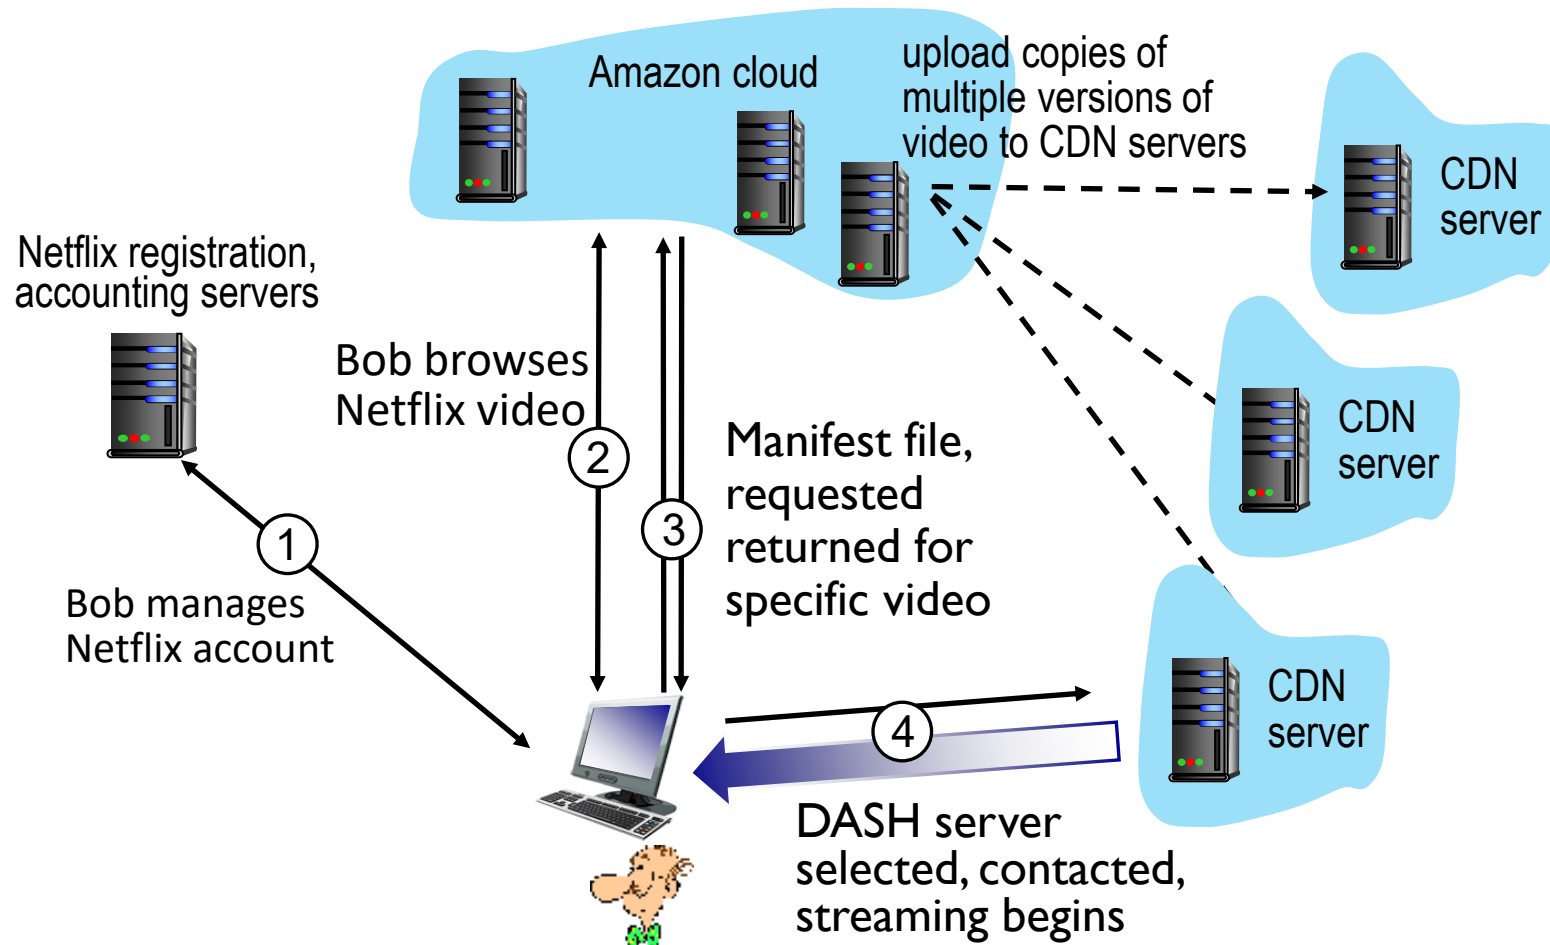


Content distribution networks (CDNs)

- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



Case study: Netflix



Chapter 2: outline

2.1 principles of network applications

- app architectures
- app requirements

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.6 Video Streaming and CDN

2.7 Socket Programming

Chapter 2: Summary

Most importantly: learned about *protocols*!

❖ typical request/reply message exchange:

- client requests info or service
- server responds with data, status code

❖ message formats:

- *headers*: fields giving info about data
- *data*: info(payload) being communicated

important themes:

- centralized vs. decentralized
- stateless vs. stateful
- scalability
- reliable vs. unreliable message transfer
- “complexity at network edge”